

xv6 Survival Guide
Introduction to Operating Systems

Mark Morrissey
markem AT pdx DOT edu

Acknowledgements

This document is intended to assist students at PSU¹ taking CS333, Introduction to Operating Systems. In particular, sections using the xv6 operating system[4] for projects involving operating systems design and implementation.

We are fortunate to have had feedback from many in the development of this guide and the approach to the class in general. In particular, we want to acknowledge the contributions from our course staff, especially Jeremiah Peschka, Ted Cooper, Joe Coleman, Enis Inan, Abbie Jones, Kaleb Striplin, and Naomi Dickerson.

There have also been students who have contributed to the course either through suggestions and code snippets at weekend help sessions, helpful links and tips sent in private email, and running experiments beyond the scope of the class. We thank Jacob Collins, Ethan Grinnell, Charlie Juszczak, Eric Laursen, Andy Mayer, Saly Hakkoum (Electrical and Computer Engineering), Philipp Wesp (exchange student from Heidelberg University) and especially Noah Zentzis for their contributions.

¹Portland State University, Computer Science Department, PO Box 751, Portland, OR 97207. [Department home page](#).

Contents

1	Introduction	1
1.1	Administrivia	1
1.2	Regrade Requests	1
1.3	D2L Usage and Deadlines	1
1.4	Special Needs & Extraordinary Circumstances	2
1.5	Programming Languages	2
1.5.1	Coding Style	2
1.5.2	The C Preprocessor	2
1.5.3	Makefile Considerations	2
1.6	Prerequisites	3
1.6.1	Processes	3
1.6.2	System Calls and Low-Level I/O	3
1.6.3	Concurrency	3
1.6.4	Algorithms and Data Structures	4
1.7	Resources	4
1.8	The Project Report	4
1.8.1	LaTeX	5
1.8.2	Project Report Structure	5
2	Preparing Your Submission	7
2.1	runoff.list	7
2.2	Creating and Testing Your Submission	7
3	About xv6	9
3.1	Getting Started	9
3.1.1	Download and Unpack xv6	9
3.1.2	Building the Kernel	9
3.1.3	Running the Kernel	10
3.2	QEMU	13
3.3	Remote Access	14
3.3.1	No XWindows (nox)	14
3.3.2	Google Chromebook	14
3.4	Debugging	14
3.4.1	Alternate Symbol Table	14
3.4.2	Debugging a Crashed Kernel	14
3.4.3	Tracing Instruction Pointers	15
3.4.4	Inspecting Key Data Structures	16
3.5	XV6 C Library	17

A	gdb Practice	19
A.1	Part 1	19
A.2	Part 2	20

Chapter 1

Introduction

Operating systems is a very large topic. In this course, you will learn about core concepts and implement some of those concepts. The purpose of this document is to give you guidance and advice that course staff feel will be beneficial. This document is **required reading**.

This document is *not* a replacement for attending lectures nor is it a crutch that will allow you to avoid assigned readings. This document assumes that you are reading assigned materials before the subject matter lecture, attend and participate in lectures, and keep current on material in the D2L discussions area.

1.1 Administrivia

This section of CS333 will, in addition to lectures, include a series of projects that will add interesting features to xv6. The lectures will explain concepts and strategies used in operating systems while the projects will reinforce these concepts and give you hands-on experience with implementing some strategies in a small operating system called xv6[4].

The main course text, OSTEP[2], is a free e-textbook, the main reference for xv6 is similarly free[4]. Other resources may be helpful as we will use the C programming language[13][14]. We will also look at some Intel assembly on occasion. You likely already have a copy of “K&R”[14].

1.2 Regrade Requests

Regrade requests must be received via email to the instructor within two weeks of the assignment grade and feedback appearing on D2L. Requests must be specific as to why the score or feedback is incorrect. A response will be provided in email. The email chain will be archived.

1.3 D2L Usage and Deadlines

Our class will use D2L for projects. You will use D2L to submit your projects. We will provide feedback and grading there as well. One of the reasons that we are using D2L is that the assignment drop box section of D2L can be set to reject any submissions beyond a certain date and time. Late assignments will be penalized 10% of the total points per hour late. No assignments will be accepted that are over two hours late. Check the syllabus for more information on due dates.

1.4 Special Needs & Extraordinary Circumstances

Accommodations are collaborative efforts between students, faculty, and the Disability Resource Center. Students with accommodations approved through the DRC are responsible for contacting the faculty member in charge of the course prior to or during the first week of the term to discuss accommodations. Students who believe they are eligible for accommodations but who have not yet obtained approval through the DRC should contact the DRC immediately. The DRC may be reached at 503-725-4150 or visit the DRC web page: <http://www.drc.pdx.edu/>.

If an extraordinary situation (for example severe illness) prevents you from working for a period of time, contact us as soon as possible to discuss your situation and arrange a special schedule. Scheduled work commitments do not constitute an extraordinary circumstance. Makeup exams will not be given except in cases of severe and documented medical or family emergencies. Please note that travel is not considered an emergency. If an emergency arises and you miss an exam, contact the instructor before the exam to arrange for a special circumstance. Medical or similar documented emergencies that fall within University guidelines for accommodation will receive accommodation. Know the University rules before you ask as the burden of meeting the guidelines falls upon you, the person requesting accommodation.

1.5 Programming Languages

Students will be exposed to assembly language code written in GNU assembly for the Intel IA-32 architecture (x86) and code in the C programming language for both kernel and user level programs.

All programming projects will be in ANSI C and we will use the GNU C compiler[8]. The xv6 kernel does contain 32-bit assembly code for the Intel x86 architecture, which you are expected to be able to understand. We will not be using advanced features of the Intel instruction set. The 1987 version of the manual is relatively small, easy to understand, and will suffice for our purposes[12].

Students should be aware that the C library provided with xv6 (`ulib.c`, `printf.c`, and `umalloc.c`) is a highly restricted subset of the standard C library. There is no written documentation for this library but students can easily learn the details by reading the source, which is only 280 lines of code. Talk to the instructor before you consider expanding the C library.

1.5.1 Coding Style

In any large project, matching the existing coding style is important. Having different coding styles intermixed leads to confusion and bugs. Students are **required** to follow the existing coding style in xv6. In particular, the pay close attention to function declarations and how the *function name* begins the *line after* the function return type. For *helper functions* which are limited in scope to a specific file, you **must** declare the function as “static”. If this is unclear, see §4.6 of [14].

1.5.2 The C Preprocessor

You have already seen the C preprocessor in action if only for the simple case of a `#include` statement near the top of your source code. Section 4.11 of [14] concerns the C preprocessor. We will use additional features of the preprocessor, especially the `#define` statement and `#ifdef ... #endif` for conditional compilation (§4.11.3). The GNU documentation for the C preprocessor[7] is excellent.

Conditional compilation will be used in this course for all project code.

1.5.3 Makefile Considerations

The Makefile for this class has the student set the project number as an integer at the top of the file. Unfortunately, if you leave in any trailing spaces, the make system misunderstands your intent. If you

are a vim user, you can set up your `.vimrc` file to automatically remove trailing spaces. Here is an email provided by the course TA in a previous section:

A number of students have run into problems with trailing whitespace in Makefiles causing conditional compilation to fail. The easiest way to fix this is to tell your editor to remove trailing whitespace.

Big red warning: I'm not a regular vim user, so if you run into problems I won't be able to troubleshoot your vim configuration and I'll send you to either the tutors or Stack Overflow (where I found this solution)

For brave vim users, here's the fix. You need to add the following to your `.vimrc` file (it's located in your home directory on the Linux systems):

```
fun! <SID>StripTrailingWhitespaces()
    let l = line(".")
    let c = col(".")
    %s/\s\+$//e
    call cursor(l, c)
endfun
autocmd FileType c,cpp,java,php,ruby,python autocmd
    \ BufWritePre <buffer> :call <SID>StripTrailingWhitespaces()
autocmd BufWritePre Makefile :call <SID>StripTrailingWhitespaces()
```

The second to last autocmd is a long one and it's probably wrapping in your mail client - make sure it's only one line in your `.vimrc`

1.6 Prerequisites

CS201, [Computer Systems Programming](#) is a prerequisites for this course. You are assumed to have taken this course and are prepared to implement certain concepts in an applied setting within xv6. The current text for CS201 is "B&O"[\[3\]](#). There should be a copy of the CS201 text in the library, if you do not have your own. The course staff encourage you to refer to this text throughout the term.

1.6.1 Processes

Chapter 3 of B&O[\[3\]](#) covers process structure and execution. This information will be quite important as we discuss operating system support for processes. In particular, sections 3.4, 3.6 and 3.7 are seen as very useful for students to review.

1.6.2 System Calls and Low-Level I/O

Chapter 8 of K&R, [The UNIX System Interface](#), discusses system calls that provide low-level I/O capabilities to programs and also discusses memory allocation and management. Chapter 10 of B&O covers related topics. If you are not already familiar with this material, please look at it and ask questions as soon as possible. Tracing and implementing system calls will be important in xv6.

1.6.3 Concurrency

The CS201 text contains an excellent chapter on concurrent programming¹. Concurrency is a very important topic in operating systems. You are encouraged to read this chapter to assist you in understanding concurrency, especially threads. Concurrency will be a major recurring topic throughout the course.

¹See [\[3\]](#) Chapter 12.

1.6.4 Algorithms and Data Structures

Lower division courses such as CS163, [Data Structures](#), CS202, [Programming Systems](#), and CS201, [Computer Systems Programming](#) cover fundamental algorithms and data structures. This course will emphasize algorithms, data structures, and their correct implementation in a new setting; namely concurrent programming. Students are assumed to be familiar with simple data structures in the C Programming Language.

1.7 Resources

The following resources are good to know about.

1. Study Groups. Students are strongly encouraged to form study groups. Be sure to go prepared to your study group so that you can participate fully and get the most out of this valuable resource. Remember to study concepts and ideas but do not share solutions.
2. D2L Discussions. You are **required** to monitor the discussion area of D2L for this course. Course staff monitor and contribute to this area. Project clarifications will be posted here. Students are responsible for all clarifications and additional information posted.
3. CS Linux Lab. You will develop your projects using department systems in the Linux lab (use babbage.cs.pdx.edu, do not use linux.cs.pdx.edu). The Linux lab systems have the required virtual machine and tool chain installed. All assignments will be graded on these machines as well. Your assignments must work correctly on these machines or they will be considered incorrect.
4. CS Tutors. The [CS tutors](#) can help you with some questions around the C programming language and the GNU debugger, `gdb`[10]. They cannot help you with programming projects.
5. Online Tutorials. There are several good online tutorials. If you have some you wish to share, let us know and we will list them here.

GDB

[CProgramming.com](#)

[Unknown Road](#)

[OSDev.org](#)

6. Class TAs and Technical Course Support Specialists. Class personnel are selected on their ability to help guide you through the rigors of the course. TAs will hold regular office hours and TCSSs will conduct labs.

1.8 The Project Report

Among the deliverables for this class is the project report. The purpose of the project report is to document your work on each project and communicate to course staff important details of your design decisions, test methodology, and key aspects of your implementation.

Your project report will be submitted with each assignment and are **required** to be in PDF format. Hand-written, simple text, and MS Word reports **will not be accepted**. Systems such as MS Word and \LaTeX are ideal for producing properly formatted PDFs. There are free packages that claim to produce documents compatible with MS Word and these include OpenOffice, FreeOffice, LibreOffice, and Google Docs. Your instructor is familiar with MS Word and \LaTeX and recommends \LaTeX for project reports.

1.8.1 LaTeX

LaTeX is a typesetting system that is very suitable for producing scientific and mathematical documents of high typographical quality. It is also suitable for producing all sorts of other documents, from simple letters to complete books. LaTeX uses TeX as its formatting engine. From [15].

You are encouraged to try out LaTeX and, to assist your efforts, the LaTeX source for this and other course documents will be provided[6]. The systems in the Linux lab have all the LaTeX tools installed but the instructor prefers to use TeXstudio on his Windows machine as it has a great GUI. The easiest path for the Windows user is to get ProTeXt from TUG (TeX Users Group). A very good reference for LaTeX is [15] and the wikibook is also good.

Students have recommended sharelatex.com. Sharelatex.com is a web application that is a resource for people who want to try LaTeX but perhaps don't want to commit or aren't sure how to install an editor. It has all the basic functionality and several nice features, including vim and emacs bindings. Best of all, they have documentation that includes tutorials and guides on most things one would want to do.

1.8.2 Project Report Structure

The reports will have the following sections, in the specified order. An example report template is provided for you under "Project 1 Example Report" on the course website. Seek clarification from the course staff when some issue is not clear.

1. Assignment Description. This section will contain, in your own words, a general sense of what the assignment required you to do and learn. It should not have too many details of what you specifically did, as those belong in the deliverables and implementation sections. Please do not merely copy the project description. Note that this section should not be overly verbose. Instead, think of it as a short, one or two paragraph summary of the key ideas in the assignment plus any new items that you've learned while doing the assignment.
2. Assignment Deliverables. This is a description of the principle artifacts that you produced for the project. In general, the assignments will require you to add new system calls, create user commands, and also modify, improve, and add to the existing functionality of the various facilities in xv6 (e.g. the scheduler, file system). These are all mentioned in this section, along with a brief description of the added functionality. So, this section should indicate what the new system call lets the user do; what the new user command does and that, if it displays something, a brief description of what it is displaying; and, finally, the general nature of the newly modified, improved, or added functionality. Note that you must list all assumptions, limitations, and edge cases [that are implementation dependent] of each feature in this section (e.g. that the displayed elapsed time is to the nearest hundredth of a second, that the maximum UID and GID is 32,676 (unsigned), etc.).
3. Implementation. This describes the actual code you wrote to implement the deliverables. When you reference it, use the specific filename and the corresponding line number(s). Do not give a step by step rundown of what the code does (e.g. I declared two integer variables, subtracted them, ...). Instead, only describe the non-obvious aspects of your code, such as a description of what a helper method does, what a block of code is intended to do (e.g. Lines X-Y in proc.c grabs a process from the free list), a layman description of the algorithm you used for some specific calculation/action, etc. You must list all of the files that were modified here, and describe the modifications made in each one. *Do not include the actual source code unless it is absolutely necessary to communicate the implementation.* For example, it is fine to include source code if you need to provide a step-by-step explanation of a particularly tricky algorithm, but only include the necessary code. Assume that the reader is familiar with xv6 and has access to your code. In general, there should be no need to list any code here.

4. Testing Methodology. This is where you will document your tests and testing strategy to show that your project meets its requirements. Do not include source code in this section except in those (rare) cases where a code snippet would greatly improve understanding. Every test must have a corresponding output demonstrating the test. You risk losing points if you fail to properly document your tests; “it should be obvious” does not suffice. In general, tests must follow this outline:

- (a) Test Name
- (b) Test Description
- (c) Expected Results
- (d) Test Output / Actual Results
- (e) Discussion
- (f) Indication of PASS/FAIL

For “Test Output”, use screen shots, not copy-paste. Some projects will provide useful test code that will let you test some, if not all, of the requirements outlined in the “Required Tests” section. You are free to use any provided test code to show that you meet the relevant requirements, but you must describe how each of the provided tests work and what requirements they specifically test in order to get any credit for that test.

Chapter 2

Preparing Your Submission

Project submissions will be compiled and graded using the Linux lab. It is the responsibility of each student to ensure that their submission compiles and runs correctly in that environment.

Students are **required** to follow the directions in this section for generating an archive file for submission. Failure to follow these directions will result in a points deduction.

2.1 runoff.list

Add any new files that you want to be submitted as part of your assignment to the file `runoff.list`. Here are example entries:

```
# CS333 additions for Mark Morrissey
date.c
ps.c
ps.h
time.c
```

Add new file names at the end of the file.

2.2 Creating and Testing Your Submission

You are **required** to test your submission to ensure that it works. You test that your code works with the command `make dist-test-nox`. This command will create the `dist` and `dist-test` subdirectories and then run the xv6 kernel for testing. If you find that files are missing then it is likely that you failed to update the `runoff.list` file. When you are satisfied with the results, create your submission with `make tar`. The tar file is a copy of your “dist” subdirectory.

The tar file created by the `make tar` command constitutes the code submission for each project. Students **may not** submit Apple “.rar” files, “.zip” files, etc. Project reports are to be submitted separately from the source code. This means that each projects submission will consist of two files.

Chapter 3

About xv6

Your project will require you to “evolve” the xv6 operating system[4]. Xv6 is designed to be small, compact, easy to understand and modify, and similar in structure to Linux and Unix systems. By the end of the course, not only will you have a bootable operating system of your own, you will also possess practical skills that will directly transfer to industry.

3.1 Getting Started

3.1.1 Download and Unpack xv6

The source code for the xv6 kernel is packaged in a file named `xv6-pdx.tar`; click the link to download the file. From the Linux lab you can also copy the source with this command:

```
cp ~markem/public.html/CS333/xv6/xv6-pdx.tar .
```

To unpack the source code from the tar file, use this command:

```
tar xf xv6-pdx.tar
```

You will now have a directory named `xv6-pdx`. Change to this directory (`cd xv6-pdx`).

List the contents of this directory:

```
$ ls
asm.h  bio.c      bootasm.S  bootmain.c  buf.h  BUGS
cat.c  console.c  cuth       date.h      defs.h
...
```

As you can see, the kernel is comprised of many parts. There are also several “helper” files¹. You will become very familiar with a small group of these files and leave many of them alone.

3.1.2 Building the Kernel

The project `Makefile` differs substantially from the original MIT version. See Project One for details on using this file. Be careful to follow the rules regarding conditional compilation.

Build the kernel using these two commands:

```
make clean
make
```

¹Any file not ending in `.c`, `.h`, or `.S` is a helper file, for our purposes.

The default “make rule” is to build the kernel. The kernel is always built in a way that let’s us use it with the GNU gdb debugger. While we would not want to make the kernel in quite this way in a production environment, it is a very good way to build the kernel during development.

There are also several targets for running the kernel in different modes. Our class will only use two of these targets. The primary make targets that we will use are:

```
qemu-nox: fs.img xv6.img
qemu-nox-gdb: fs.img xv6.img .gdbinit
```

There are others that we will not use in this class. Note that a target that allows for execution within gdb includes “-gdb” and one that does not use an Xterm includes “-nox”. We will use “qemu-nox” or “qemu-nox-gdb” when running our kernel.

3.1.3 Running the Kernel

Since we are in a development environment using a simulated machine (via QEMU), the way that we run our code is a little bit more complex than just turning on a computer and watching the monitor. Fortunately, our method is also very flexible and does not require dedicated hardware.

Normal Boot

When you “boot” the kernel normally, you will use the command

```
make qemu-nox
```

This will build the kernel, if necessary, then load and run xv6. You will know that your kernel has properly loaded and is executing the shell when you see output similar to

```
qemu-system-i386 -nographic -hdb fs.img xv6.img -smp 2 -m 512
xv6...
cpu1: starting
cpu0: starting
sb: size 2000 nblocks 1941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$
```

The “\$” is the shell prompt. The other lines are output when xv6 is initializing various parts of the kernel. We will not look at initialization in depth in this class.

At this point, you have a running operating system but one that can do very little. We will discuss the limitations and expand the capabilities through the course, but for now just type “ls” and you will see the entire list of shell commands available to you. The list is very small. Also note that the xv6 shell does not understand the concept of a search path, so you need to either be in the directory with the command you wish to execute or give the full path to the command.

Initializing with GDB

Remote access² will be via some SSL-based tool such as putty or slogin. You will need a minimum of two remote logins; one for running the kernel and one for the gdb debugger session³.

Once connected and in the correct directory, use the command

```
make qemu-nox-gdb
```

²If you are logged in at a console in the Linux lab, just use two separate shell windows.

³If you use a tool such as **screen** or **tmux** then you can just use sub-windows inside that tool.

This will build the kernel, load it into the qemu simulator and wait for gdb to connect. Sample output looks like this⁴:

```
$ make qemu-nox-gdb
gcc -Werror -Wall -o mkfs mkfs.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -fvar-tracking -fvar-tracking-assignments -
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -fvar-tracking -fvar-tracking-assignments -
gcc -m32 -gdwarf-2 -Wa,-divide -c -o usys.o usys.S

[. . . many lines of the compilation process deleted . . .]

dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes (512 B) copied, 0.0171023 s, 29.9 kB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
353+1 records in
353+1 records out
181117 bytes (181 kB) copied, 0.00223519 s, 81.0 MB/s
sed "s/localhost:1234/localhost:27377/" < .gdbinit.tmpl > .gdbinit
*** Now run 'gdb'.
qemu-system-i386 -nographic -hdb fs.img xv6.img -smp 2 -m 512 -S -gdb tcp::27377
```

The last line states that the boot of the kernel is waiting for a gdb session to connect. You do this by typing “gdb” from a login session in the same directory as you ran the make command. The first time, you will see an error message like this:

```
$ gdb
GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.2) 7.7.1
Copyright (C) 2014 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
warning: File "/u/markem/xv6-pdx/.gdbinit" auto-loading has been declined by your 'auto-load safe-
To enable execution of this file add
    add-auto-load-safe-path /u/markem/xv6-pdx/.gdbinit
line to your configuration file "/u/markem/.gdbinit".
To completely disable this security protection add
    set auto-load safe-path /
```

⁴Example output text may run off the right edge of the page. The output that is truncated is just an example. You will be able to see the entire lines of output when you run the commands for yourself.

line to your configuration file `"/u/markem/.gdbinit"`.
For more information about this security protection see the
"Auto-loading safe path" section in the GDB manual. E.g., run from the shell:
 `info "(gdb)Auto-loading safe path"`
(gdb)

It is a bit wordy, but the part to pay careful attention to is this:

To enable execution of this file add
`add-auto-load-safe-path /u/markem/xv6-pdx/.gdbinit`
line to your configuration file (e.g., `"/u/markem/.gdbinit"`).

What this *really means* is to edit the file `“.gdbinit”` in your home directory (you may need to create it) and add the line that corresponds to the one in bold above to enable gdb to connect to your kernel. Note that your error message will, of course, be slightly different. Once you have added this command, you can exit gdb and then reenter and you should see something like this:

```
$ gdb
GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.2) 7.7.1
Copyright (C) 2014 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
+ target remote localhost:27377
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i8086 settings.

The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: jmp $0xf000,$0xe05b
0x0000fff0 in ?? ()
+ symbol-file kernel
(gdb)
```

Finding Our Initial Breakpoint

Now that we have the kernel built, find the address of the `_start` variable in the kernel.

```
$ nm kernel | grep _start
8010b50c D _binary_entryother_start
8010b4e0 D _binary_initcode_start
0010000c T _start
```


In this case, the address is 0010000c.

Running the Kernel

Start QEMU running with the xv6 kernel and the gdb debugging “hook” enabled. You should have two windows open to the Linux lab and the current directory for both windows should be the project directory. We will designate the one running xv6 the “command window” and the one running gdb the “debug window”. In the command window, enter the command

```
make qemu-nox-gdb
```

This will start the QEMU simulator and pause the loading of the kernel at the start of the boot code, waiting for gdb to connect.

Attaching gdb

In the debug window, launch gdb; you should see output similar to the following

```
$ gdb
GNU gdb 6.8-debian
Copyright (C) 2008 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.  Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
+ target remote localhost:26000
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: jmp $0xf000,$0xe05b
0x0000fff0 in ?? ()
+ symbol-file kernel
(gdb)
```

You can now use gdb normally.

3.2 QEMU

Note: the version of QEMU installed on CAT machines has a bug. This bug is harmless for our purposes but does occasionally print a spurious message: main-loop: WARNING: I/O thread spun for 1000 iterations. This message can safely be ignored.

You will “boot” and run your operating system using QEMU[16]. QEMU is a powerful emulator offering very good performance. You will not be expected to delve into the QEMU layer of the project. The project is structured so that QEMU “looks” like a native X86 processing environment. The principle advantage of QEMU is that dedicated hardware is not necessary for your machine so that kernel crashes and hangs are recovered from easily. This environment will greatly speed your learning and also shorten development cycles.

3.3 Remote Access

There are several ways to run xv6 from a remote login to the Linux lab. The preferred way is to use the “-nox” option that runs the qemu[16] environment without using XWindows. There is a noticeable speed improvement when using -nox. Additionally, there are certain issues with the XWindows server for Apple computers that are avoided using this technique.

3.3.1 No XWindows (nox)

By default, QEMU will create a console window that is an xterminal. To suppress this behavior and run your console in the same window from which you launch QEMU, use `make qemu-nox` or `make qemu-nox-gdb`. To exit from the console, you will need to use one of two QEMU commands:

1. control-a x. This will terminate the QEMU console and so terminate xv6.
2. control-a c. This will drop you into the QEMU console⁵ where you can use “quit” to exit.

3.3.2 Google Chromebook

To use a Chromebook, there are two steps:

1. Install the secure shell extension from the Google store. A new shell is created with control-shift-N.
2. Follow the instructions for No XWindows (§3.3.1).

3.4 Debugging

At times, you will be running your operating system from within the GNU debugger, gdb[10]. The xv6 environment includes a hook for allowing gdb to debug a running kernel. You should have some familiarity with gdb, but if you get stuck and the documentation is not illuminating, ask a specific question on the course mail list or see the CS tutors.

Note that the first time you run the kernel under gdb, you are likely to get an error message about not being able to access the local .gdbinit file. In order to allow gdb to properly connect to your kernel, you will need to modify (or create) the .gdbinit file in your home directory to include a specific auto-load command, such as

```
add-auto-load-safe-path /u/markem/CS333/pdx-xv6/.gdbinit
```

Note that the gdb error message will specify the correct information to put in the .gdbinit file.

3.4.1 Alternate Symbol Table

By default, the xv6 environment will load the symbol table for the executing kernel. However, there are times when you will want to debug user programs inside of xv6. The gdb manual in section §2.1.1 **Choosing Files** covers selection of different symbol tables from other files. Inside of gdb, just use `symbols file` to select a new symbol table file. This new file is usually the .o file created during compilation of your user-level program.

See also §16 **Examining the Symbol Table** for additional gdb commands related to symbol tables.

3.4.2 Debugging a Crashed Kernel

The xv6 kernel “crashes” via a call to `panic()`. Unfortunately, this does not cause the kernel to halt, but rather “hang”. In your gdb session, you can enter control-c to break to the gdb command prompt and inspect state from the gdb prompt. The `panic()` function will print out instruction pointer⁶ information

⁵Further information regarding console commands for QEMU can be found in section 3.5 of [16].

⁶Called program counter (PC) in xv6.

and you can view `kernel.asm` in an editor to trace the call path that (hopefully) led to the panic.

3.4.3 Tracing Instruction Pointers

The terms “instruction pointer” and “program counter” are equivalent. As Intel uses “instruction pointer” (e.g., register `eip`), we will mostly use this term in lectures. Be aware that both terms have the same meaning.

When the kernel “panics”⁷, the output will look something like

```
panic: Example panic message
```

```
80104ebb 801038f7 80103896 0 0 0 0 0 0
```

Normally, there will be a number to the right of the text for the `panic` message, but there is not one here. The line of numbers, however, represents the instruction pointer trace leading to the panic. We can trace the kernel execution by looking at the instruction pointers, most recent first, while viewing the decompiled kernel, `kernel.asm`. We must remember that the instruction pointer will point to the *next* instruction to be executed.

In `kernel.asm`, we search for address `80104ebb`. We find this code in the region of the instruction pointer:

```
p = ptable.pReadyList[i];
80104e97:      8b 45 f4      mov    -0xc(%ebp),%eax
80104e9a:      05 cc 08 00 00  add    $0x8cc,%eax
80104e9f:      8b 04 85 a4 39 11 80  mov    -0x7feec65c(,%eax,4),%eax
80104ea6:      89 45 f0      mov    %eax,-0x10(%ebp)
if (p) {
80104ea9:      83 7d f0 00      cmpl   $0x0,-0x10(%ebp)
80104ead:      74 0c          je     80104ebb <scheduler+0x95>
panic("Example_panic_message\n");
80104eaf:      c7 04 24 2f 91 10 80  movl   $0x8010912f,(%esp)
80104eb6:      e8 7f b6 ff ff      call  8010053a <panic>
// Enable interrupts on this processor.
sti();

// Loop over free list looking for process to run.
acquire(&ptable.lock);
for (i=0; i<NPRIQ; i++) {
80104ebb:      83 45 f4 01      addl   $0x1,-0xc(%ebp)
80104ebf:      83 7d f4 02      cmpl   $0x2,-0xc(%ebp)
80104ec3:      7e 81          jle    80104e46 <scheduler+0x20>
```

which we would recognize as being part of the scheduler. This shows us the call to `panic()`. Don’t worry, with practice you will get the hang of reading `.asm` files.

Now take a look at the address `801038f7`:

```
xchg(&cpu->started, 1); // tell startothers() we're up
801038d7:      65 a1 00 00 00 00  mov    %gs:0x0,%eax
801038dd:      05 a8 00 00 00      add    $0xa8,%eax
801038e2:      c7 44 24 04 01 00 00  movl   $0x1,0x4(%esp)
801038e9:      00
```

⁷Makes a call to the `panic()` function.

```
801038ea:      89 04 24          mov    %eax, (%esp)
801038ed:      e8 df fe ff ff      call   801037d1 <xchg>
scheduler();    // start running processes
801038f2:      e8 2f 15 00 00      call   80104e26 <scheduler>

801038f7 <startothers>:
```

which we can see corresponds to a call to our `scheduler()` routine.

Finally, we look at the last address `80103896`

```
userinit();    // first user process
8010388c:      e8 4b 0f 00 00      call   801047dc <userinit>
// Finish setting up this processor in mpmain.
mpmain();
80103891:      e8 1a 00 00 00      call   801038b0 <mpmain>
80103896 <mpenter>:
}
```

which shows that the chain started with a call to `mpmain()`.

This instruction pointer trace (a chain of instruction pointers) shows that the kernel called the routine `mpmain` then called the routine `scheduler()` and then called the `panic()` routine. If we look at the file `main.c` at the end of the `main()` function we will see the call to `mpmain` that starts this all off. With practice, you should have a pretty good handle on tracing instruction pointers.

3.4.4 Inspecting Key Data Structures

The most common data structures that we will investigate are associated with processes. The `ptable` and `proc` structures are used to manage processes.

The `ptable` struct defined in `proc.c`:

```
struct {
    struct spinlock lock;
    struct proc proc[NPROC];
} ptable;
```

which contains an array of `proc` structures. This struct is defined in `proc.h`:

```
enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };

// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t *pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;    // Process state
    uint pid;               // Process ID
    struct proc *parent;     // Parent process
    struct trapframe *tf;    // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;       // Current directory
    char name[16];          // Process name (debugging)
};
```

If we are running the kernel from inside gdb and we get a panic message such as:

```
cpu1: panic: panic for printing out lock state
```

```
80104e6a 801038f7 801038b0 705a 0 0 0 0 0
```

Then we can go to the gdb session and enter control-c:

```
(gdb) c
Continuing.
^C
Program received signal SIGINT, Interrupt.
The target architecture is assumed to be i386
=> 0x8010555c <xchg>:  push  %ebp
xchg (addr=0x801139a0 <ptable>, newval=newval@entry=1) at x86.h:122
122      {
(gdb)
```

Since I happen to know where the `panic()` message is at, I decide that I want to track the instruction pointer values associated with the lock in the `ptable` structure which is defined in `proc.c`. Here is a listing from the “pcs” (program counters) array in `ptable.lock.pcs`:

```
(gdb) p/x ptable.lock.pcs
$1 = {0x80104e3d, 0x801038f7, 0x801038b0, 0x705a, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0}
```

which shows the most recent program counter first.

To trace these values, use the same procedure as in §3.4.3.

3.5 XV6 C Library

It is likely that students are used to using certain functions from the C library on a system such as Linux. It is important to note that the C library for xv6 is very minimal. Much of the library is implemented in `ulib.c` but some functions are implemented in separate files. For example, the `printf()` function for xv6 is implemented in `printf.c` and provides only a minimal subset of the functionality that students may be used to. The full list of routines in the C library can be found in the file `user.h`. In a later project, you will need to modify the `ls` command and its implementation in `ls.c`.

While it is possible to expand the functionality of xv6 C library routines, and even add new ones, note that this is *not* the purpose of this class. All assignments can be completed without modifying or extending the existing C library beyond what is described in the assignments. An obstacle for many students is be the lack of formatting options in `printf()` but careful thought will allow you to get formatting as necessary although some coding will be required.

As far as **documentation for the xv6 C library**, there isn’t any. The xv6 book does describe some facets of some library calls but mostly it is left to the programmer to read the code to determine proper usage. This is normal in systems development and good to practice.

Appendix A

gdb Practice

This chapter walks the student through using the `gdb` debugger to trace and inspect the boot process for the `xv6` kernel. This task will help familiarize students with key techniques that have proven useful for debugging in the `xv6` environment.

Follow the instructions to boot `xv6` into the debugger given in §3.1. Once you have `xv6` booted into the debugger and are at the debugger prompt, continue.

A.1 Part 1

Set a `gdb` breakpoint at the location of `_start` and then continue to run the debugger until it arrives at that location in memory.

```
(gdb) br * 0x0010000c
Breakpoint 1 at 0x10000c
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x10000c: mov    %cr4,%eax

Breakpoint 1, 0x0010000c in ?? ()
(gdb)
```

The exact details of the output may vary slightly. Look at the registers and the stack contents:

```
(gdb) info reg
...
(gdb) x/24x $esp
...
(gdb)
```

Determine which part of the output is contained on the stack; not all the listing is part of the stack. Try to determine what the non-zero parts of the stack correspond to. That is, why are the values there, how did they get there, and what part of the code put the values there. You might find it convenient to consult the files `bootasm.S`, `bootmain.c`, `bootblock.asm` and `kernel.asm`. The reference page has pointers to `x86` assembly documentation, if you are wondering about the semantics of a particular instruction. Here are some questions to help you along. Do not answer these questions in your homework, they are just here to help you if you are stuck.

- Start by setting a break-point at 0x7c00, the start of the boot block (bootasm.S).
- Single step through the instructions (type `si` to the gdb prompt). Where in bootasm.S is the stack pointer initialized? Single step through the call to bootmain; what is on the stack now?
- What do the first assembly instructions of bootmain do to the stack? Look for bootmain in bootblock.asm.
- Look in bootmain in bootblock.asm for the call that changes `eip` to 0x10000c. What does that call do to the stack?

A.2 Part 2

In bootasm.S, find this line:

```
ljmp    $(SEG_KCODE<<3), $start32
```

This line moves our processor into 32-bit mode. The reason that this happens is that, when the code was laid out, `start32` was declared to be a 32-bit *segment* with the declaration `.code32`. The change to 32-bit addressing occurs automatically because the processor “sees” that we are jumping to a new segment (the long jump). Once in the new 32-bit segment, the code sets up known values for registers, updates the stack pointer (`%esp`) to point to `$start` and then calls the routine `bootmain` which resides in `bootmain.c`.

Now look at bootblock.asm around line 100. You should see this line:

```
call    bootmain
7c48:    e8 e2 00 00 00    call    7d2f <bootmain>
```

The first line is the assembly code for the call to bootmain and the second line is that code in machine language. The value `7c48` is the memory address for the call to bootmain. The right-hand side of the second line shows us that bootmain is at memory address `7d2f`.

Now let’s look further into the bootblock.asm file to find address `7d2f` and we see that it is indeed the entry to bootmain:

```
00007d2f <bootmain>:
```

This routine is checking to see if what is stored “on disk” is a file in the expected form (an ELF binary) and then loads data from the file (the data will be executable instructions). Then there is the call to `stosb`, which is very important. The ELF header contains information as to where the start location (entry point) for this binary is located.

We can see that the entry point is extracted and turned into a function call with this code:

```
entry = (void (*)(void))(elf->entry);
entry();
```

The resulting machine code is one line:

```
7db1:    ff 15 18 00 01 00    call    *0x10018
```

and we can tell that the function that will be called has its memory address stored at location `0x10018`. Inspect this address to see the value stored there. What does this tell you about how the memory location for bootmain gets set? Hint: it is set long before you began to boot the kernel.

Now, open the file `main.c`. Compare the first few lines of the routine `main` with the output from `gdb` when you give the `list` or `li` command.

Bibliography

- [1] Alex Aiken, *MOSS: A System for Detecting Software Plagiarism*, <http://theory.stanford.edu/aiken/-moss/>
- [2] Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau, *Operating Systems: Three Easy Pieces*, Arpaci-Dusseau Books, V0.9, May 2015.
- [3] Randal E. Bryant and David R. O'Hallaron, *Computer Systems: A Programmer's Perspective, third edition*, Pearson Education, 2016, ISBN 978-0-13-409266-9. [Website](#).
- [4] Russ Cox, Frans Kaashoek, and Robert Morris, *xv6: a simple, Unix-like teaching operating system*, Draft as of September 3, 2014, xv6-book@pdos.csail.mit.edu. [Local copy](#).
- [5] Russ Cox, Frans Kaashoek, and Robert Morris, *xv6: a simple, Unix-like teaching operating system*, source code listing. [Local copy](#).
- [6] Mark Morrissey, *xv6 Survival Guide*, [PDF](#), [L^AT_EX source](#).
- [7] Free Software Foundation, Inc., *The C Preprocessor*.
- [8] Free Software Foundation, Inc., *GCC Version 5.3.0*.
- [9] Free Software Foundation, Inc., *GCC Make Manual*.
- [10] Free Software Foundation, Inc., *Debugging with gdb: the gnu Source-Level Debugger (GDB)*.
- [11] [Git User Manual](#).
- [12] Intel Corp., *Intel 80386 Programmer's Reference Manual*, 1987, [website](#).
- [13] Al Kelley and Ira Pohl, *A Book on C, Fourth Edition*, Addison Wesley, 1998, ISBN 0-201-18399-4.
- [14] Brian W. Kernighan and Dennis M. Ritchie, *The C Programming Language, Second Edition*, Prentice Hall, 1988, ISBN 0-13-110362-8.
- [15] Tobias Oetiker, Hubert Partl, Irene Hyna and Elisabeth Schlegl, *The Not So Short Introduction to L^AT_EX 2_ε, Version 5.05*, July 18, 2015.
- [16] QEMU: Open Source Processor Emulator. [QEMU Documentation](#).
- [17] W. Richard Stevens and Stephen A. Rago, *Advanced Programming in the UNIX Environment, Third Edition*, Pearson Education, 2013, ISBN 978-0-321-63773-4.
- [18] MIT course [6.828: Operating System Engineering](#).